

SITE PLAN & CONTENT SPECIFICATION · WORKING DRAFT V2

A Framework for Ethical and Constructive use of AI in Early Childhood Education

The build spec for a field-standards website: what each page contains, what ships in what order, and the draft substance for the pages that carry the argument.

Prepared for Ben Clemens Date 10 September 2026 Posture Neutral field resource
Readers Providers · Intermediaries

CONTENTS

THE SITE

01 Design constraints

The field spans a wider capacity range than almost any other sector this framework could be written for, and the two ends fail differently. A single rule set will fail one of them.

The small end, which is most of the field

The modal early childhood provider is a single site where one person is the director, the compliance officer, the HR department, and the IT department — or a home-based provider who is the entire

01	Design constraints
02	Sitemap
03	Home page
04	Start here: four questions
05	Principles
06	Where AI earns its place
07	Bright lines
08	What the law says
09	Policies to adopt
10	The multi-site track
11	Cases
12	Vendors
13	FAQ

THE BUILD

organization. Median hourly wage for an early educator is **\$13.07**. Roughly 30% of center-based staff and 20% of home-based providers hold a bachelor's degree. Only about 37% of even public-school pre-K teachers have had any training in evaluating an ed-tech product's quality.

1. **Assume no legal review.** Anything that says "have counsel review your vendor agreement" fails for most of the field. Ship completed language, not instructions to draft it.
2. **Assume the reader is the whole organization.** No RACI charts, no "convene your AI governance committee." One person is reading this between nap time and pickup.
3. **Assume twenty minutes, not twenty hours.** A one-page policy a provider can read, sign, and date gets adopted. A forty-page framework gets downloaded once.
4. **Spanish at launch, not later.** A third of children ages 0–5 are dual language learners. A framework with a principle about language access that ships English-only undercuts itself on day one.

The large end, which nobody is writing for

A twelve-site chain, a large single facility, or a Head Start grantee with delegate agencies has an HR director, a compliance officer, and access to counsel. Those four rules are wrong for them — and being left out is why they default to one of two bad answers: a blanket ban that pushes AI onto

personal devices where nobody can see it, or an unmanaged accumulation of tools no two sites agree on. Section 10 is their track, and it needs its own rules:

5. **Give them a governance structure, not a longer policy.** They do not need a template; they need delegated approval authority, a pilot protocol, and a defensible answer for a board.
6. **Treat the reflexive no as a failure, and say so.** A ban that staff quietly work around leaves the same exposure with less visibility. This has to be named directly, because the organizations that need to hear it believe they have already solved the problem.
7. **Use their leverage on behalf of everyone else.** Large buyers can require contract terms a single center can only ask about. Model rider language published here, adopted by enough large operators, changes vendor defaults for the whole field.

Both ends produce the same structural rule, which shows up throughout this spec: **tier by provider capacity as well as by risk**. The same underlying guidance becomes a one-page signed document for a family child care home, a six-page policy for a center, and a governance charter with contract riders for a multi-site operator. The principles and the bright lines do not change across those three. Everything else does.

NAMING The working title is descriptive, which is

right for a document and wrong for a domain. The site needs a short name before launch — **Birth to Five AI Standards** is the strongest candidate: it uses the field's own age framing rather than "early childhood," which readers outside ECE hear as kindergarten. Avoid anything built on "First Five," which collides with California's First 5 commissions.

02

Sitemap

Eleven destinations. The order is deliberate: the front door is a tool, not a manifesto, because a director arrives with a specific product in front of her and a decision to make this week.

SECTION	WHAT LIVES THERE
Home /	The intake tool as the main action; three bright lines, an FAQ preview, the artifact shelf. Spec in Section 1.
Start here /start	The four-question intake. Return to risk tier, the guidance for that tier, the right policy template for that provider type.
Principles /principles	Seven named principles, short and quotable, dated and versioned. Every other page cites them by name.
Where AI helps /good-uses	Twelve constructive uses with vignettes, tiers, and guardrails — operations, revenue, and funding.

	that returns hours to children. See 06.
Bright lines /bright-lines	Eight prohibited uses, each with legal or evidentiary basis. Deliberate short. The page most likely to be screenshotted.
What the law says /law	The ECE regulatory crosswalk — setting by setting, then a state law. Does not exist anywhere else.
Policies to adopt /policies	Sixteen artifacts in editable form with fill-in markers, last-updated stamps, English and Spanish.
For larger operators /multi-site	Governance charter, delegated authority, model contract riders, protocol, multi-state mapping, shadow-AI discovery. Section 10.
Cases /cases	Worked vignettes — a real situation, what's at stake, what to do, which principle governs.
Vendors /vendors	Ten questions, an answer-record sheet, the disclosure card format, eventually the register.
For intermediaries /intermediaries	Separate track: how a CDFI, R&R T&TA provider assesses AI in its lending, underwriting, and coaching — plus co-branding guidance.
Policy recommendations /recommendations	Asks aimed at state legislatures, Office of Head Start, state licensure and QRIS. Kept distinct from pro

	guidance.
FAQ /faq	Why this exists and why the obvious alternatives don't cover it. Carried over from the previous argument so no other page has to cover it. See Section 13.
About & changelog /about	Members, advisers, funders, met with them to find out what we don't know yet, how the current own AI use worked, and a dated changelog.

NOTE Keep [/recommendations](#) genuinely separate. A site that mixes "here is what your program should do" with "here is what the legislature should pass" reads as advocacy to the provider and as vague to the legislator.

03 Home page

Six modules, in this order. No large hero image, no mission statement above the fold — a director who has landed here has a question, and the page's job is to route her in one click.

HOME · MODULE STACK (6)
1 · ROUTE Thinking about an AI tool? Answer four questions. One sentence of framing, then the intake. Not a carousel, not a video. The four questions can render

inline on the home page itself rather than behind a click — they are short enough, and every extra step loses the reader whose next task is lunch.

Start

2 • WHERE IT HELPS

Three places this is already saving programs real money.

Three of the twelve constructive uses, each a one-line scenario with the hours or dollars it returns. This module comes *before* the bright lines deliberately: a site that opens with prohibitions gets read as hostile by the director who is drowning, and she is the reader.

3 • BRIGHT LINES

Three things not to do, whatever the tool promises.

Three of the eight, rotated or fixed, each one line with a link to the basis. This is the module most likely to be screenshotted and forwarded, which makes it the site's cheapest distribution.

4 • THE SHELF

Take what you need.

Direct download links to the four artifacts that do the most work — the one-page policy, the staff one-pager, the family notice, and the vendor question sheet — with the Spanish version beside each. No form, no email capture. Anything that gates a policy template behind a signup will lose the providers who need it most.

5 • FAQ PREVIEW

Why does this exist?

Four questions surfaced as collapsed rows,

expanding in place, with a link to the full FAQ. These four are chosen to answer the objections a director or a funder forms in the first ten seconds — and they are where the case for the site lives, so no other page has to carry it.

Aren't states already legislating on AI in early childhood?

+

Doesn't NAEYC already cover this?

+

Can't we just use the K-12 AI guidance?

+

We're one classroom. Isn't this for big organizations?

+

6 • STANDING

Who wrote this, and when.

Member names, funders, a version number, and a last-reviewed date — visible on the home page, not buried in About. In this posture legitimacy is the product, and a date stamp is the cheapest possible proof that the guidance is maintained.

ONE RULE FOR THE WHOLE PAGE

Nothing above the fold explains what AI is. The reader already knows; what she doesn't know is whether she is allowed to use the thing her vendor just switched on.

04

Start here: four questions

A mad-lib rather than a questionnaire. The user completes one sentence from dropdowns; the guidance behind it stays locked until the sentence is done. Four toggles do the entire tiering job, and every one is answerable without knowing anything about AI.

THE SENTENCE

We are a [family child care home / center / multi-site or Head Start grantee / intermediary] considering a tool that [does / does not] see or receive information naming an individual child or family, [does / does not] capture a child's face, voice, or body, [will / will not] inform a decision about a specific child, and [will / will not] be used directly by children.

The first dropdown is the one that matters most for adoption: it tiers by provider type as well as by risk, so the same answer set returns a one-page document for a home-based provider and a six-page one for a grantee.

What it returns

TIER 1

Back office, no child or family data

Marketing copy, menus, newsletters naming no child, professional development, general

	grant boilerplate, staff scheduling against ratio requirements. Returns: acceptable-use one-pager, disclosure when AI drafted family-facing text, human review before anything is sent.
TIER 2	Operations touching child or family data Subsidy billing, enrollment records, licensing documentation, CACFP reconciliation, grant narratives containing child vignettes, family communication and translation. Returns: written agreement required, register entry, no consumer chatbots, retention and deletion terms, disclosure log where Head Start applies.
TIER 3	About an individual child, or used by a child Screening and assessment support, observation narratives, photo and video portfolios, behavior tracking, referral routing, classroom cameras, anything a child interacts with. Returns: full vendor due diligence, separate parental consent, DLL-disaggregated evidence, documented human determination, and an exit plan before adoption.

TIER 0

Do not adopt

The eight bright lines. Run as a screener *before* the tiering questions, so a program that lands here never sees implementation guidance at all.

WATCH "Back office is safe" is the one shortcut this framework must not take. In early childhood the back office *is* where the child data lives — subsidy billing contains family income, immigration-adjacent documentation, and custody information; grant narratives routinely contain child vignettes. Tier 1 has to be defined by the absence of child and family data, not by the word "administrative."

05

Principles

Seven, named as one word plus a claim in the imperative so they survive being quoted in a board meeting. Two of these exist in no other AI framework anywhere, and those two are why this site is not a link to the K-12 guidance.

RELATIONSHIP

Use AI to return adult attention to children.

Development from birth to five is built through serve-and-return interaction with a consistent, attuned adult. Anything that

occupies the adult's attention or substitutes for the interaction hits the mechanism of development directly. This inverts the usual efficiency argument into a pedagogical one: the strongest case for AI on paperwork is that paperwork is time stolen from children. Zero to Three's appropriate-use recommendation for educators is exactly this — administrative tasks and planning, supporting adults' time rather than replacing human caregivers.

DEVELOPMENT

Judge every tool against what a child this age actually needs.

No other AI framework has a developmental-appropriateness test, because every one of them was written for K-12 or for generic nonprofits. RAND found the dominant hesitation among pre-K teachers is developmental appropriateness and screen time — not privacy. Meeting the field where its actual concern lives is both honest and strategically smart. *Drafting note: the AAP replaced its prescriptive time limits with the "5 C's" framework in January 2026 — do not restate the 2016 hour limits, which most ECE materials still cite.*

STANDING

Decide for children who cannot decide for themselves.

A four-year-old cannot notice an AI system, object to it, recognize it as non-human, or report harm from it, and leaves no record of dissent. This is qualitatively different from K-12, where student voice is at least available in principle. Parental consent is the only consent available and it is structurally compromised: signed at enrollment as a condition of care, often by families with no alternative provider, sometimes under subsidy conditions, frequently in a second language. The National Center for Youth Law's *digital legacy* framing belongs here — data created about a three-year-old in 2026 may follow them for sixty years.

CUSTODY

Know where a child's record goes, and be able to get it back.

Head Start's [45 CFR 1303.24](#) already requires disclosure logs. [§1303.23\(e\)](#) gives parents an *unqualified* right to review any written agreement the program has with a third party — the text is not limited to agreements naming their child, which makes it the strongest existing hook for a parent asking to see an AI vendor

parent asking to see an AI vendor contract. IDEA gives parents a right to have records destroyed once no longer needed — a right FERPA alone does not confer, and one an AI vendor that has ingested IFSP text into a retained embedding store cannot honor. And the amended COPPA Rule requires separate verifiable parental consent for disclosing children's personal information to a *third party* for AI development, which the FTC has said is not "integral" to a service and so cannot ride on general consent. Note the gap: the Rule is quieter about an operator training its own models on data it already holds. That is a contract term to write, not a law to cite.

CONSEQUENCE

No AI output decides anything about a child.

This generalizes rules Head Start programs already live under.

§1302.33 prohibits using screening or assessment to exclude children from enrollment, or to rank, compare, or evaluate individual children for rewards or sanctions. §1302.17 prohibits expulsion for behavior.

Generalized: no AI output may be the basis for excluding,

disenrolling, suspending, or expelling a child, or for a special education referral, without an independent human determination on the underlying evidence. Anchoring the guardrail in existing compliance obligations rather than novel ethics is what makes it adoptable.

LANGUAGE

Serve the third of children who are learning two languages.

There are 7.1 million dual language learners ages 0–5 — a third of all young children, and more than 20% of the young-child population in over half of all states. They are not a special case; they are the population. Every AI failure mode compounds for them: speech recognition performs worse on accented and child speech, screening instruments require cultural and linguistic adaptation, translated consent forms are where informed consent quietly fails. §1302.102(c)(2) already requires DLL-disaggregated analysis, so the framework can require DLL-disaggregated performance evidence from vendors as an extension of an existing duty rather than a new ask.

CURRENCY

Say when this was written and what has changed since.

The field's flagship technology position statement is fourteen years old and predates the smartphone-saturated classroom. Meanwhile Colorado's AI Act was delayed, blocked in federal court, and replaced by a different statute inside eighteen months. Every artifact carries a date, a version, and a changelog entry. Guidance that cannot say when it was last reviewed should not be trusted, and saying so in your own principles is what earns the right to say it about the 2012 statement.

CUT CANDIDATE If seven feels long, **Currency** is the one to demote — it governs the framework rather than the use of AI, and could live as a masthead commitment instead. Six is probably the sweet spot.

06 Where AI earns its place

A site that is only a list of prohibitions will be read once and never returned to. Most of what AI can do in early childhood is unglamorous, low-risk, and worth real money to a business running on 1–3% margins — and every hour it returns is an hour an

adult spends with children.

This is the Relationship principle made concrete, and it is also the page that makes the bright lines credible. A framework that only says no reads as hostile to a director who is drowning; a framework that says *here is where this genuinely helps, and here is the short list where it doesn't* reads as useful. Each entry below carries its risk tier and its guardrail, because the point is not that these uses are unregulated — it is that they are worth doing carefully.

FOR PROVIDERS

The corrective action plan TIER 2

A licensing visit turns up three findings — a gap in the medication log, an expired CPR certification, and playground surfacing depth. The corrective action plan is due in ten business days. The director has done this before. The fixing was the easy part; writing it up took most of a weekend.

RETURNS	The remedies still take what they take. The document describing them drops from a weekend to an evening. Across a year — licensing renewal, two monitoring visits, a policy manual refresh — this is the difference between a director who does paperwork on Sundays and one who doesn't.
GUARDRAIL	The findings and the remedies are hers; only the drafting is the tool's. Corrective action plans often reference specific

incidents, so use an approved tool with an agreement on file and strip child identifiers first.

Subsidy billing that actually reconciles

TIER 2

A center serves 34 subsidized children across two counties. Attendance sheets, authorization certificates, and payment remittances never quite line up. Every month a handful of days go unbilled or get paid at the wrong rate, and chasing each one costs more staff time than it recovers — so nobody chases them.

RETURNS	This is not efficiency; it is revenue. On margins this thin, recovering a few hundred dollars a month of correctly billed care is the difference between a raise for a teacher and no raise. The reconciliation produces a discrepancy list with the underlying evidence attached, in minutes.
GUARDRAIL	The most sensitive data in the building — family income, immigration-adjacent documentation, custody arrangements. Written agreement, register entry, no consumer tools. The tool flags; a person decides what gets submitted.

Finding the money you were always eligible for

TIER 1

A family child care provider has been in business eleven years. In that time she has learned about exactly two funding sources: CACFP, because a peer mentioned it, and a state stabilization grant, because it arrived in the mail. She has never heard of the facility improvement

fund her county runs, and she has never completed an application with a narrative section.

RETURNS	A profile-matched list of what she is actually eligible for — QRIS quality incentives, CACFP, county and First 5 funds, SBA microloans, USDA rural facility programs — with deadlines and required documents. Then a first draft of the narrative in her own words, from a ten-minute conversation instead of a blank page. The application she would never have started gets submitted.
GUARDRAIL	Verify every opportunity, amount, and deadline against the source before applying — a hallucinated deadline costs a provider real money. And the moment a narrative describes an individual child, the work moves to Tier 2.

The empty slot

TIER 2

A center runs at 87% enrollment. Two infant spots have been open five weeks. Nine families sit on the waitlist, but nobody has called them, because the person who would make the calls is covering ratio in the toddler room at three o'clock.

RETURNS	Unfilled slots are the largest single revenue leak in the sector, and how fast an inquiry gets answered largely determines whether it becomes a tour. An assistant that responds within minutes, works the waitlist, and books tours converts a structural leak into filled seats — two infant slots at typical rates is meaningful annual revenue against a razor-thin margin.
GUARDRAIL	Families must know they are talking to

GUARDRAIL

Families must know they are talking to an AI, and a person takes over the moment the conversation touches a child's needs, a disability, or a subsidy question. California SB 243-style chatbot duties may attach. It never makes an enrollment decision — that is a Consequence-principle line.

Every family gets the notice in their own language

TIER 2

A program serves families speaking Spanish, Vietnamese, Mam, and Arabic. Routine communication — the closure notice, the menu, the reminder about water bottles — goes out in English, because there is no realistic alternative. Families who don't read English hear about the snow closure from a neighbor, or don't.

RETURNS

Routine operational communication in every home language, instantly, at no marginal cost. This is probably the single most equity-positive low-risk use of AI in the sector: it fixes a daily failure that nobody counts as a failure, for a population that is a third of all children ages 0–5.

GUARDRAIL

Routine only. Screening results, IFSP and IEP conversations, behavior conversations, eligibility determinations, and anything touching suspected maltreatment go to a qualified interpreter — bright line 6, and §1302.33 already requires it. Where a bilingual staff member or family can review the routine translations, have them do it; Indigenous languages like

Mam are exactly where machine translation is weakest.

Knowing what a child actually costs you

TIER 1

A provider sets her rates by looking at what the center down the street charges. She has never calculated true cost per child by age group, so she does not know that her infant room loses money on every enrollment and her preschool room quietly subsidizes it. When she considers expanding, she is guessing.

RETURNS	A working cost model built from her own numbers through a conversation, rather than a spreadsheet she would never finish. Break-even by classroom, the real cost of a staffing change, what a rate increase does to enrollment risk. This is a business-health intervention — and it is precisely what an intermediary's business coaching is trying to accomplish, arriving between coaching visits instead of once a year.
GUARDRAIL	No child or family data, so Tier 1 — but it is her own financials, which means a tool with a real agreement rather than a free consumer account. Treat the output as a draft for her coach to check, not an answer.

Staffing to ratio without overstaffing

TIER 1

Ratios are legally binding and shift by age group, by hour, and by who called in sick. The director builds next week's schedule on paper on Sunday night, over-staffs to be safe, and absorbs the labor cost.

RETURNS	A schedule that holds ratio through the
---------	---

day's actual arrival and departure curve. Overstaffing is a margin killer; understaffing is a citation. The gap between them is where the money is — and, incidentally, where the Sunday nights are.

GUARDRAIL

Staff schedules are employment data. The tool proposes; the director decides. Do not let scheduling optimization drift into performance monitoring of a low-wage workforce that did not agree to be measured this way.

The CDA portfolio that stalled two years ago

TIER 1

A teacher earning fourteen dollars an hour is working toward her Child Development Associate credential. The competency statements and the professional portfolio are the part that stalls people — not the knowledge, the writing.

RETURNS

A study partner that quizzes her, helps her turn what she already does every day into the reflective statements the portfolio asks for, and is available at nine at night after her shift. Credential attainment raises wages and reduces turnover, which is the sector's most expensive unsolved problem.

GUARDRAIL

The reflections have to be hers — a portfolio assesses her practice, and submitting generated text is both dishonest and useless to her. The tool asks the questions; she answers them. No child names in the examples.

FOR THE NONPROFITS AND INTERMEDIARIES THAT SUPPORT THEM

What the coaching notes have been saying all along

TIER 2

A resource and referral agency holds six years of technical assistance notes across 400 providers — thousands of pages of free text that nobody has ever read as a set. Everyone in the building has a theory about where providers get stuck. Nobody has evidence.

RETURNS	Qualitative coding at a scale that was never affordable: which topics recur, which cluster with later closure, where the coaching model is repeating itself, and what providers keep asking for that the organization does not offer. An archive becomes a program design input — and the kind of finding a funder will pay to act on.
GUARDRAIL	Notes carry provider and sometimes family identities. De-identify before analysis, or work in an approved environment with an agreement in place. Findings describe patterns across a portfolio; they never quietly become a score attached to an individual provider.

The assessment write-up TIER 2

A business coach spends half a day on site conducting a Business Administration Scale assessment, then two more days writing the report. The assessment is the expertise. The report is transcription.

RETURNS	write-up time falls by roughly two-thirds, which converts directly into more assessments per coach per year — the same staff reaching more providers. The coach's scores and judgments are the input; the tool assembles the document around them.
GUARDRAIL	Every score and every recommendation is the assessor's. The tool never generates a rating, and a person reads the finished report before it reaches the provider. Reports name real businesses and real financials, so this is approved-tool, agreement-on-file work.

A funding radar for the whole portfolio

TIER 1

A CDFI knows about opportunities its 300 portfolio providers don't. It has never had a way to match them one to one, so it sends a monthly newsletter that most providers do not open.

RETURNS	Matched alerts — this provider, this opportunity, this deadline, these documents — instead of a broadcast. The intermediary's information advantage finally reaches the individual provider it would help, and application rates rise because the ask is specific rather than general.
GUARDRAIL	Match on business characteristics the provider has actually shared, never on inferred ones. Verify every opportunity before it goes out: a wrong deadline costs a provider money and costs the intermediary the trust that makes it useful.

Where to send the coach next

TIER 3

A network supports 400 providers with eight coaches. Triage today is whoever called, whoever a coach happened to visit, and whoever a funder asked about.

RETURNS	A ranked view of where technical assistance would do the most good, built from signals the organization already holds — enrollment trends, subsidy billing patterns, licensing history, time since last contact. Eight coaches deployed against need rather than against the phone.
GUARDRAIL	This is the closest thing on the page to a bright line and it needs the strongest one. It may route <i>support</i> — never sanctions, funding denials, or portfolio exits. Every provider can see what the system says about them and contest it. And the model must be tested for whether it systematically ranks home-based providers, newer providers, or providers in child care deserts as less worth helping. The thin-file problem is real here, and it points the wrong way.

THE RULE THAT KEEPS THIS PAGE HONEST

Every entry on the live page ends with what the program actually did with the time — hired a floater, gave a teacher a paid planning hour, opened an infant slot, stopped working Sundays. Without that, this is an efficiency brochure, and the field has plenty of those. With it, it is the Relationship

principle with receipts. Collect those outcomes from real pilots rather than borrowing vendor case studies, name no products, and carry the tier and the guardrail on every single one.

07 Bright lines

A short prohibited-use list is usable; a five-by-five impact-likelihood grid is not. Each line pairs a rule with the legal or evidentiary basis that makes it defensible, because a director will have to justify it to an owner or a board.

Emotion, affect, or engagement recognition on children.

EU AI Act Art. 5(1)(f) prohibits emotion inference in education institutions outright — and it reaches placing on the market and putting into service, not just end use, so it binds vendors too. The mechanism is documented: Halberstadt et al. (2020) found prospective teachers 1.36 times more likely to show racialized anger bias toward Black children's faces. Human raters carry that bias, training labels carry it, models learn it. There is no U.S. analogue to the EU prohibition — which is precisely why the field should write one.

AI as the basis for expelling, suspending, disenrolling, or denying enrollment to a child.

Already law for Head Start under §1302.17 and §1302.33 . Generalizing it costs nothing and matters

enormously: three- and four-year-olds are expelled from state-funded preschool at roughly three times the K-12 rate, and from private and community programs at more than thirteen times. Black children account for nearly half of public preschool suspensions while representing under a fifth of enrollment. *Vintage caution: the multipliers trace to Gilliam's 2005 work and the suspension disparity to the 2013–14 Civil Rights Data Collection. Pull the most recent CRDC cycle before publishing.*

AI as the basis for a special education or early intervention referral, without an independent human determination on the underlying evidence.

The failure mode is fully documented in adjacent child-serving systems. Oregon discontinued its child-welfare screening tool after it flagged a disproportionate number of Black children; the Allegheny Family Screening Tool drew a DOJ civil-rights inquiry over disability bias. Referral is a consequential decision about a child and carries the strongest human-decision requirement in the framework.

AI anywhere in mandated-reporter documentation.

No authority anywhere addresses this — no state child welfare agency guidance, no CAPTA guidance, nothing from any ECE body. The unresolved questions are real: when does reasonable suspicion attach if a tool flags first? Is an AI-smoothed narrative still the reporter's contemporaneous observation? A flat prohibition is simple, defensible, and costs providers nothing. This is a place the site can be first.

Child or family personal information into a consumer chatbot with no written agreement.

Under §1303.22 a Head Start grantee that pastes child observations into a consumer tool has arguably made an undisclosed third-party disclosure — no consent, no reviewable agreement, and no entry in the §1303.24 disclosure log. The violation is complete whether or not anything leaks, which is the point providers most often miss.

Machine translation in IFSP or IEP meetings, screening-result conversations, child-protection conversations, or eligibility determinations.

Routine operational communication — newsletters, reminders, menus — is fine and genuinely useful. High-stakes content is where machine translation fails, with documented safety consequences in medicine and explicit cautions from the National Center for State Courts. §1302.33 already requires qualified bilingual staff or interpreters for exactly these conversations.

Face or voice biometric identification of children — check-in/out, photo auto-tagging — without express, separate, revocable parental consent.

Illinois BIPA applies to "private entities" and expressly excludes state and local government agencies, which flips the risk map most people assume: a private center or a nonprofit grantee carries more exposure here than a public school district does, with a private right of action and liquidated damages. The amended COPPA Rule expressly includes biometric identifiers, including voiceprints and facial templates, in "personal information." Consent must come from a parent on behalf of a child who cannot give it — a live legal question with no clean answer.

Vendor training of models on children's images, voices, work, or records without separate

verifiable parental consent.

Not aspiration — current federal law, with one boundary to state precisely. The COPPA Final Rule effective June 2025, full compliance April 2026, requires separate verifiable parental consent for disclosing children's personal information *to third parties* for AI development; the FTC treats such disclosures as not "integral" to a service and so outside general consent. It says less about an operator training its own models on data it already holds — so this line is enforced by contract as well as by statute. Photo and video portfolio apps are the highest-volume collectors of children's images, video, and voice in the sector.

08

What the law says

This is the single most valuable page on the site, and no one has written it. Every vendor contract template in circulation cites the wrong statutes for early childhood.

The page answers one question — *which law actually governs the data in my program?* — by setting, because the answer changes completely depending on who operates the program.

SETTING	FERPA	WHAT GOVERNS INSTEAD
State pre-K run by a school district	Yes	FERPA, plus state student-vendor law (e.g. CA Educ. C 2-d)

Head Start that is itself a FERPA-covered educational agency Typically district-operated, but the test is the funding, not the org chart	Yes	45 CFR 1303.21(a) direct funds under a Department and therefore subject to FERPA <i>instead of</i> the Head
Head Start run by a CAA, nonprofit, tribe, or private entity	No	45 CFR 1303 Subpart C by disclosure exceptions, parental amendment, right to review agreements, disclosure log §1303.21(b) displaces Sub B/C confidentiality for any eligible under IDEA, in any
Private or community child care center	No	COPPA (on the vendor), state consumer privacy law, license own privacy policy and family
Family child care home	No	Same as above, with effect privacy floor at all
Community-based Part C early intervention	Not directly	IDEA Part C confidentiality broader "participating agency identifiers covered, parental

The three findings that change practice

- **COPPA is the law that actually reaches every ECE setting**, because it regulates operators of services directed to children under 13 and does not care whether the program is a school — though note it binds the vendor, not the provider, so a program's leverage runs through its contract. And the FTC explicitly declined to codify the ed-tech "school authorization" exception in the 2025 rule, to avoid conflicting with anticipated FERPA amendments. So the widely-relied-upon practice of a program consenting on parents' behalf remains policy-statement guidance built for schools, which most ECE providers are not. That is a genuine, citable legal uncertainty the site can be the first to name.
- **California's early-learning privacy law is not SOPIPA.** AB 2799 (2016) did not amend SOPIPA; it added a separate statute at B&P §22586 — the Early Learning Personal Information Protection Act — covering operators of services used and marketed primarily for preschool and pre-K purposes, with biometric information expressly included in "covered information." An ECE vendor in California may be regulated by ELPIPA and not SOPIPA, and templates citing SOPIPA/AB 1584 are citing the wrong law.
- **HIPAA usually does not apply, and people get this backwards.** A program that does not bill electronically for health services is not a covered entity; immunization records and

screening results in a child's file are not PHI in that provider's hands. The constraint runs on the inbound side — what the program may re-disclose is governed by the authorization it signed plus Head Start or IDEA. Edge cases worth flagging: Early Head Start embedded in an FQHC, programs billing Medicaid for screening.

The state layer

Built as a lookup rather than prose, and honest about churn. Priority states by DLL and provider concentration: California, Texas, New York, Florida, New Jersey, Illinois. Each entry answers four things — does the state's student-privacy law reach pre-K or child care; is there a biometric statute with a private right of action; what do the state's CCDF plan and provider agreement require; and what has the state AI office or education department published that a provider might be handed.

OPEN QUESTION Two state readings need a lawyer before publication. New York Ed Law 2-d defines "school" to include publicly funded pre-K, which on its face pulls community-based organizations delivering publicly funded pre-K inside the statute — a big deal for New York City providers. And whether Illinois SOPPA reaches preschool at all could not be confirmed. Both should be owned by a pro bono legal partner.

09

Policies to adopt

Sixteen artifacts, not one policy. Every one ships editable, with [SAMPLE] markers where local text goes, a last-updated stamp, and a Spanish version. Publish two variants of the core policy rather than one canonical version — it lowers the stakes of choosing.

FAMILY CHILD CARE · SINGLE SITE The One-Page AI Policy Printable, plain language, sign-and-date. What we use AI for, what we never use it for, what happens to your child's information, who to ask. One side of one sheet. If only one artifact gets built well, it is this one.	CENTER · MULTI-SITE Program AI Use Policy Four to six pages. Governance and approval, data handling by tier, approved tools, staff obligations, family notice, incident response, annual review commitment. Structured so a director can delete sections that don't apply.
HEAD START GRANTEE Part 1303 Addendum Maps AI use onto existing obligations: third-party disclosure under §1303.22 , the disclosure log under §1303.24 , parent review of written agreements under §1303.23 , and screening constraints under §1302.33 . A	ALL STAFF Staff Acceptable-Use One-Pager Posted in the staff room. Green list, red list, and one sentence on what to do when unsure. Assumes the reader has thirty seconds and no context.

compliance document,
not an ethics document.

FAMILIES

Family Notice

What AI the program
uses, in what part of the
work, and what
information it touches.
Plain language at a
middle-school reading
level. Kept strictly
separate from the
consent form — notice
is not consent, and
conflating them is how
programs get into
trouble.

FAMILIES

**Consent: Images,
Voice & Model
Training**

The COPPA-driven
instrument. Separate,
specific, revocable
consent for biometric
capture and for any
vendor use of a child's
images, voice, work, or
records in model
development. Includes
a plain-language
explanation of what
"training" means,
because most families
will not know.

DIRECTOR

**Approved Tools
Register**

Fillable. Tool, purpose,
tier, who approved it,
date, agreement on file,
data it touches, review
date. The artifact most
likely to be requested in
a licensing visit or a
monitoring review.

DIRECTOR ·
INTERMEDIARY

**Vendor Question
Sheet**

Ten questions with a
place to record the
answer and the date.
Providers have roughly
zero negotiating
leverage with a major
vendor — what they
can do is ask,
document, and choose.
Section 12.

BOARD · OWNER

Board Talking Points

One page a director can

ANY PROGRAM WITH A
HANDBOOK

Handbook Addendum

hand to a board or owner to get a policy approved without having to teach AI first.

Language
Amend what exists rather than write something new. Drop-in paragraphs for an existing family handbook, staff handbook, and photo-release form. The lowest-friction path to adoption by far.

DIRECTOR

Incident & Discontinuation Protocol

What to do when a tool fails, leaks, or turns out to do something it didn't disclose: who to notify, what to tell families, how to stop using it, and what to document. Nobody has written the provider-side version of this.

MULTI-SITE · CHAIN · GRANTEE

Governance Charter & Authority Matrix

Who approves what, at which tier, without asking whom. The artifact that prevents both failure modes in Section 10 — the queue that staff route around, and the accumulation nobody inventoried.

MULTI-SITE · CHAIN · GRANTEE

Model Contract Rider

Drop-in vendor terms for buyers with actual leverage: no training on our data including internally, no biometric processing without written instruction, deletion reaching derived artifacts, subprocessor notice,

MULTI-SITE · CHAIN · GRANTEE

Pilot Protocol & Stop Condition

One site, a named owner, a fixed window, success criteria declared in advance, staff and family notice, and a written condition under which the pilot ends. Most enterprise pilots lack the last one,

DLL-disaggregated validation on request, notification with a clock. The highest-leverage document on the site.

MULTI-SITE · CHAIN · GRANTEE

Shadow-AI Amnesty Kit

A two-week, no-consequence window for staff to register what they are already using, with the announcement language and the intake form. Run before writing policy — otherwise the policy governs a fiction.

which is why they never end.

CDFI · R&R · T&TA

Intermediary AI Practice Standard

The other track. How an intermediary handles AI in its own lending, underwriting, portfolio monitoring, and coaching — plus co-branding guidance for delivering the provider artifacts under its own name.

THE MECHANIC THAT COMPOUNDS

Ask programs to send back the policy they actually adopted. Within two years that makes this body the only one that knows what the field's AI policies really say — a research asset, a revision input, and a legitimacy claim at once.

10

The multi-site track

The design constraint that governs the rest of this site — no lawyer, no IT department, twenty minutes — is right for the modal provider and wrong for a fourteen-site organization with an HR director and outside counsel. Larger operators are not a smaller problem. They are a different one, and the field has no guidance for them at all.

A large center, a chain, or a Head Start grantee with delegate agencies has something the single site does not: capacity to actually govern this. What it does not have is anyone telling it how. Left alone, these organizations fail in two opposite directions, and both are worse than they look.

Failure mode one: the reflexive no

Risk-averse counsel or a cautious board issues a blanket prohibition. It feels like the safe choice and it is documented as a policy, so everyone moves on.

What actually happens is that staff use AI anyway — on personal phones, in free consumer accounts, with child names in the prompt, outside any register, any agreement, and any disclosure log. The organization now has the same exposure it had before, minus the ability to see it, plus a written policy proving it knew better. **A blanket ban is not a safety posture. It is an unmonitored one.** This should be said plainly on the page, because the people who need to hear it are exactly the ones who believe they have solved the

problem.

The corollary matters too: a program that forbids AI outright is choosing to leave the Section 06 savings on the table — the billing recovery, the translated closure notices, the grant it never applied for — while its staff take the risks anyway.

Failure mode two: the unmanaged yes

Scale cuts the other way just as easily. One site pilots a camera vendor. Another turns on an observation app's AI feature because it was on by default. Procurement buys an enterprise chatbot for the central office. Nobody has an inventory, no two sites are operating to the same standard, and governance debt compounds faster than at a single site because there are more people able to say yes.

What the multi-site track adds

Nine things, on top of everything in Section 09. Each exists because it does not make sense below a certain size.

1. **A governance charter with real delegated authority.** Site directors approve Tier 1 without asking. A named role approves Tier 2. A small group approves Tier 3 in writing, with the crosswalk in front of them. Without delegation everything queues behind one person, and staff route around the queue — which is how the unmanaged yes starts.

2. **One tool register, many sites.** The single-site register does not scale. One source of truth, per-site adoption status, and a standing rule that no site adopts a Tier 2 or Tier 3 tool that is not on it.
3. **Contract riders — because you actually have leverage.** This is the biggest asymmetry in the sector and the most valuable thing large operators can do. A single center can only ask the ten questions. A fourteen-site organization, or a grantee with delegates, can *require* terms: no training on our data including internally, no biometric processing without our written instruction, deletion that reaches derived artifacts, subprocessor notice, DLL-disaggregated validation evidence on request, and a notification commitment with a clock on it. Publish model rider language. If the field's larger buyers converge on the same terms, vendor defaults shift for everyone — including for the family child care provider with no leverage at all. Frame it that way: this is a public good only the big operators can supply.
4. **A pilot protocol with a stop condition.** Large organizations should pilot, not roll out. One site, a named owner, a fixed window, success criteria declared in advance, staff and family notice, and a real condition under which it ends. Most enterprise AI pilots have no stop condition, which is precisely why they never stop.
5. **Multi-state compliance mapping.** A chain operating in Illinois, California, and New York

faces BIPA, ELPIPA, and Ed Law 2-d simultaneously. Running different practices per state is possible and almost always more expensive than it looks; for most operators the cheapest defensible path is a single standard set at the strictest state's level. The crosswalk's state layer matters more to this reader than to any other.

6. **Shadow-AI discovery, without punishing anyone.** Before writing policy, find out what is already in use — a two-week amnesty window where anyone can add a tool to the register with no consequence. Skip this and the policy governs a fiction.
7. **Board and owner reporting on a cadence.** A standing quarterly item: tools in use by tier, incidents, what changed, what is pending approval. For Head Start this maps onto existing Policy Council and governing body reporting rather than adding a new obligation.
8. **Role-differentiated training.** A teacher, a family services worker, and a billing clerk face completely different risks from the same tool. Large organizations can afford to train them differently; small ones cannot. This is one of the few places where scale is a genuine advantage rather than a complication.
9. **A named escalation path.** Who does a teacher call at four in the afternoon when a tool did something strange? At a single site it is the director. At fourteen sites it has to be a named role, or it defaults to nobody.

HEAD START GRANTEEES WITH DELEGATE AGENCIES

A specific sub-case worth its own page module. The grantee is accountable for delegate compliance under Part 1303, but delegates make their own operating decisions and may adopt tools the grantee never reviewed. The register, the rider, and the escalation path all have to reach across that boundary — and the grantee's monitoring instrument should ask about AI tools the way it asks about anything else it is answerable for.

DESIGN NOTE Do not build this as a longer version of the small-provider material. The multi-site reader is not looking for a policy template; she is looking for a governance structure, model contract language, and a defensible answer for her board. Different artifacts, different page, same principles and the same bright lines.

11 Cases

Almost every comparable framework is governance — who approves, what's disclosed, how data is stored. Very little addresses the actual moment of use. Short worked vignettes fill that hole, are cheap to produce, and are the content

most likely to get forwarded between directors.

Consistent structure throughout: *the situation, what's at stake, what to do, which principle governs.*

Your observation and documentation app adds an AI feature that drafts learning stories from classroom photos. It's on by default.

AT STAKE	Whether the vendor is using children's images to train a model — and whether auto-tagging computes face geometry, which triggers biometric statutes in several states. Also whether the drafted narrative is still the teacher's observation.
DO	Turn it off until you've asked questions 3, 4, and 7 on the vendor sheet. Re-consent families before turning it on, using the separate images-and-training form rather than the existing photo release.
GOVERNS	Custody · Standing

A teacher pastes an incident description — with the child's name — into a consumer chatbot to write it up more professionally before it goes in the file.

AT STAKE	In a Head Start program this is arguably an undisclosed third-party disclosure with no consent, no reviewable agreement, and no log entry — complete whether or not anything leaks. In any program it likely violates the family contract and the program's own privacy policy.
DO	De-identify before drafting, or use an approved tool with an agreement on file. Add it to the register. If the incident touches suspected maltreatment, no AI at

all.

GOVERNS

Custody · Bright lines 4 and 5

A vendor offers your center AI camera analytics that "detects policy violations" and "understands the whole room." The pitch is about licensing compliance and liability.

AT STAKE

Continuous video of children who cannot consent, simultaneous surveillance of a low-wage workforce that is 40–56% women of color, and outputs that can feed licensing complaints, employment action, or expulsion decisions. If any component infers affect or engagement, it is a bright line.

DO

Separate the security question from the analytics question — they are usually sold together and are not the same purchase. Run full Tier 3 due diligence, and require in writing that no affect, emotion, or engagement inference runs on children.

GOVERNS

Bright line 1 · Standing · Consequence

A family child care provider uses her phone's translation app to explain a developmental screening result to a family that speaks Mixteco.

AT STAKE

The most sympathetic scenario in the library, and the guidance still has to hold. Machine translation fails on high-stakes content, and it fails worst on languages with the least training data — which is exactly the family least likely to have any alternative.

DO

Use the state's interpreter line or the

R&R's language access resources for the conversation itself. Machine translation is appropriate for the reminder that schedules it. The site should name the actual resources by state, because "use an interpreter" without a phone number is not guidance.

GOVERNS

Language · Bright line 6

Your Head Start data system ships an agentic feature that drafts family engagement notes and flags children for follow-up.

AT STAKE

Flagging children for follow-up is a decision about a specific child, arriving through a compliance system nobody thinks of as an AI tool. This is the most likely real-world path by which AI enters Head Start programs, and it will arrive as a product update rather than a purchase decision.

DO

Treat a vendor feature release as an adoption event: register it, tier it, and require the human determination before any flag informs a referral or a service decision. Ask what the flagging model was validated on and whether performance was analyzed by home language and disability status.

GOVERNS

Consequence · Language · Currency

A CDFI is offered an AI underwriting model to speed up child care facility lending.

AT STAKE

The intermediary track's own case. Credit decisions about providers who are disproportionately women of color running

	micro-businesses, using data that encodes the sector's existing inequities — home-based providers, newer providers, and providers in child care deserts are exactly the thin-file borrowers a model will treat worst.
DO	Fair lending review before pilot, adverse-action explainability, disparate-impact testing on the actual portfolio, and a human decision on every denial. Publish the practice standard so the providers being underwritten can see it.
GOVERNS	Consequence · Intermediary standard

Three more to draft for launch: **a parent asks whether her child's photos train AI** (the script a director needs, and the honest answer is often "I don't know, and here's how I'll find out"); **a grant writer runs a narrative containing child vignettes through an AI tool**; and **an AI toy arrives in a child's backpack** — for which PIRG's *Trouble in Toyland 2025* findings are the most vivid citable evidence available.

12 Vendors

Write the questions as questions, not as criteria. A single-site center has essentially no leverage with a major vendor. What it can do is ask, record the answer with a date, and choose between vendors — and a field where five thousand directors ask

the same ten questions creates leverage no individual director has.

1. **Is any information we enter, or that your product collects about children, used to train or improve your models — or anyone else's? If yes, under what consent?**
2. **Does your product compute face geometry, voiceprints, or any other biometric identifier from children?** Including for photo tagging, check-in, or search.
3. **Does your product infer emotion, affect, engagement, or behavioral risk about a child?** Ask directly; it is often buried under "insights."
4. **If a family asks us to delete their child's records, what exactly gets deleted — and does that include anything derived from those records?** IDEA gives parents a destruction right that most model architectures cannot honor.
5. **Who are your subprocessors, and where is children's data stored and processed?**
6. **Will you sign an agreement naming us as the data controller, with your role limited to processing on our instructions?**
7. **What is your written retention policy for children's personal information, and where is it published?** Now a COPPA requirement.
8. **For anything producing a judgment about a child: what was it validated on, and was performance analyzed separately for dual language learners and children with**

disabilities?

9. What happens to our data if we leave, and can we export it in a usable format?
10. Who do we call when it produces something wrong about a child, and what is your notification commitment if something goes wrong on your end?

The disclosure card, and eventually the register

A standardized one-page vendor disclosure answering those ten questions in a fixed format, published by the vendor, plus a public register of tools that have filed one. The model exists in healthcare and nowhere else; ECE directors currently evaluate vendors with nothing.

This is achievable at small scale — twenty vendors covers most of the market — and it would immediately become the site's most-linked asset. It also creates the only real incentive structure available: filing is voluntary, but a director comparing two products will notice which one didn't.

POSITIVE PATTERN WORTH NAMING

Not every finding is a warning. Cognoa's Canvas Dx — the FDA-authorized autism diagnostic aid for ages 18–72 months — returns *indeterminate* in about 37% of cases by design rather than forcing an answer. A tool that refuses to answer when uncertain

is a pattern the site should hold up explicitly. The same study's authors could not analyze performance by race or socioeconomic status because those data weren't routinely collected — which is exactly the gap question 8 exists to close.

13 FAQ

This page carries the argument for the site's existence so no other page has to. Every answer is a real objection someone will raise in the first meeting. Four of these — marked below — surface as the home page preview module.

Aren't states already legislating on AI in early childhood? [HOME](#) —

No. Not one has. FutureEd tracks 77 AI-in-education bills across 27 states this session; MultiState counts 134 across 31. Neither tracker contains a single pre-K, preschool, child care, or Head Start provision. State education-department AI guidance is now widespread and it is entirely K-12-framed.

What is moving fast is state law on chatbots, biometrics, and minors' data — 14 chatbot laws enacted in the 2026 term alone — and all of it is written for adolescents. Child safety is also the one lane where state AI law is explicitly protected from federal preemption efforts, which makes a framework built on child-safety and child-privacy obligations the most

durable ground available.

Doesn't NAEYC already cover this?

HOME

+

Can't we just use the K-12 AI guidance?

HOME

+

We're one classroom. Isn't this for big organizations?

HOME

+

Is anyone in early childhood actually using AI yet?

+

Our vendor says they're FERPA compliant and SOPIPA compliant. Isn't that enough?

+

Isn't back-office AI safe? We're not putting it in classrooms.

+

Why bright lines instead of a risk assessment framework?

+

Our legal team said no AI. Isn't that the safest choice?

+

We run fourteen sites and have a compliance team. Isn't this written for someone smaller?

+

Are you against AI in early childhood?

+

Who wrote this, and why should we trust it?

+

Is this legal advice?

+

14

Governance

In a neutral-field-resource posture, legitimacy is the product. The pattern that works: a named chair with standing outside the technology conversation, a cross-constituency body, a research partner, open funders, and a published sequence — consensus principles, then operationalized standards, then research.

Who has to be at the table

CONSTITUENCY	WHY THEY'RE INDISPENSABLE
Family child care providers & center directors	Seat the people the technology is aimed at. Without them the artifact will be the wrong length and reading level, and not useful.
Head Start grantee leadership	The only ECE setting with a real existing infrastructure. This is the one where vendor AI is arriving first.
CCR&R / R&R network	The delivery channel. They already have the compliance TA and touch thousands of providers.
State child care administrators	The path from framework to required rules, plans, provider agreements, licensing through legislation.
NAEYC, Zero to Three, Institute for Child Success	The field's existing authority. NAEYC's current statement is the gap, and a framework would be the field's answer overnight.

A privacy and civil rights lawyer	The crosswalk cannot ship without state readings that need an opinion,
Parents and family voice	Children 0–5 cannot participate, wh family-facing artifacts need familie
Vendors	The hard call. Vendors at the table g disclosure card viable; they also giv available attack. Whatever you choo prominently.
Researchers	The evidence base has real holes. A "we don't know" from a weakness in

Two credibility moves that cost almost nothing

- **Publish what you don't know.** A visible "open questions" page reads as rigor, not weakness — and it is the difference between an unfinished site and a process.
- **Disclose the site's own AI use.** For a body telling programs to disclose, doing it first is nearly free credibility, and refusing to would be noticed.

15

Build sequence

Ordered so the first release requires no convening at all. The convening is what makes phase two citable; it should not be what blocks phase one

from existing.

1

NO CONVENING REQUIRED

The operational layer

Research plus drafting plus one legal review. Also the part nobody else has built — so it stands alone even if nothing follows.

- The regulatory crosswalk, federal layer, with a lawyer's sign-off
- Bright lines, published as a working draft with bases cited
- The one-page policy and the staff one-pager, English and Spanish
- The vendor question sheet and answer-recording form
- Six of the twelve constructive uses, with guardrails — this is what keeps a launch from reading as anti-AI, and it costs nothing but writing
- The model contract rider, which large operators can start

using before anything else on the site exists

- The FAQ — it is cheap, it is mostly written, and it is what turns a thin site into an argument
- An honest "what we don't know" page

2

REQUIRES A NAMED BODY

Standing and practice

- Principles ratified and signed, with members and funders named
- The four-question intake, returning tier plus the right template
- The full policy family, including the Head Start addendum and handbook language
- The multi-site track — governance charter, pilot protocol, shadow-AI kit, delegate-agency module
- The remaining constructive uses, with outcomes collected from real pilots rather than vendor case

studies

- The case library — launch with eight or nine vignettes, add monthly
- State layer for the six priority states
- The policy-recommendations track aimed at states, OHS, licensing, and QRIS

3

REQUIRES ADOPTION

Infrastructure

- The vendor disclosure card and public register
- Generation: intake produces an editable, program-specific policy plus a matching family notice
- The intermediary curriculum, co-brandable and coach-deliverable
- The returned-policy corpus, and the first analysis of what the field's policies actually say
- Changelog discipline and an annual review

and an annual review
commitment, honored
visibly

SEQUENCING JUDGMENT

If you can only build one thing to take into partner conversations, build the **crosswalk plus bright lines plus the FAQ**. They demonstrate the field-specific expertise nobody else has, they are what a NAEYC or an Office of Head Start would react to, and together they make the case for the convening better than a deck about the convening would.

16 Decisions for you

Q1 Who convenes?

LIIF has the ECE portfolio and provider relationships but is a CDFI, not a standards body — a strength for delivery and a question mark for neutrality. Is LIIF the convener, a founding member, or the delivery partner?

Q2 Do vendors sit on the body?

It determines whether the disclosure card is viable and whether the framework reaches providers

through the products they already use. It also determines the first line of every critique. There is no neutral answer, only a disclosed one.

Q3 Bright lines as prohibitions, or as "requires board-level approval"?

Prohibitions are usable and quotable. Some partners will want softer language, especially on cameras, where programs have liability reasons for wanting them. My read: hold the prohibitions and lose the partners who need them softened — a framework that permits everything with paperwork is what the field already has.

Q4 Does the site take a position on specific products?

The register implies it eventually will, at least by omission. That is a meaningfully different institution — and a different legal posture — than one that publishes questions and lets directors ask them.

Q5 What is the Spanish-language commitment, exactly?

Simultaneous release, or translated later? Given the Language principle, "later" undercuts the whole thing. A resourcing question disguised as a production question.

Q6 Who owns legal review, and can it be pro bono?

The crosswalk is the site's crown jewel and cannot ship on inference. A law school clinic or a firm's pro bono practice is the natural fit, and having one

named on the page is itself a legitimacy asset.

17 Sources & verification

Everything asserted here traces to a source. Two categories matter differently: the durable evidence base, and the items with a short shelf life that must be re-checked immediately before anything goes public.

Verify before publishing

- **The COPPA AI-consent boundary.** The separate-consent requirement carries an express carve-out for disclosures "integral to the nature of the service," and the FTC's AI statement addresses disclosures *to third parties*. Do not write that AI training "always" requires separate consent — a vendor's counsel will catch it, and the internal-training gap is real.
- **Colorado.** SB 24-205 was delayed to June 2026, then enforcement was blocked in federal court in April 2026, then replaced by SB 26-189 in May 2026 with a notice-based regime effective January 2027. Source is a single law-firm alert. Check the enrolled bill and the docket.
- **AAP screen-time guidance.** The 2026 "5 C's" framework replaced the prescriptive hour

limits, but this rests on secondary coverage. Confirm specifically what replaced the under-24-month recommendation before writing anything about infants.

- **Head Start rulemaking.** An NPRM published May 2026 proposes removing wage-comparability and benefits requirements; records and confidentiality provisions were not proposed for change. Whether a final rule has issued is unconfirmed.
- **Expulsion data vintage.** The multipliers rest on Gilliam's 2005 work and the suspension disparity on the 2013–14 CRDC. Pull the newest CRDC cycle.
- **Product claims.** Several AI feature descriptions come from vendor marketing. Notably, brightwheel's public learning page described no AI features as of this research. Confirm directly with any vendor before naming them on a public page.
- **The 2016 Yale eye-tracking study is routinely mischaracterized.** It found teachers gazed more at Black children, especially boys — but on the vignette task there was *no* racial disparity in suspension or expulsion recommendations. Cite CRDC and Gilliam's 2005 and 2014 work for disparities; cite the 2016 study only for the attention finding. Overclaiming here is the easiest way to be discredited in front of an audience that knows this literature.

Evidence gaps to state rather than fill

No national data on broadband, devices, or IT support among ECE providers, especially home-based. No current verified national turnover figure. No meta-analysis quantifying differential screening-instrument accuracy by race or language. No verified U.S. emotion-recognition product in an ECE setting — the category is arriving through general school-security AI. No survey of AI use among Head Start grantees outside districts, private centers, or family child care homes; the RAND figures cover public-school pre-K only.

Core sources

FEDERAL LAW

[45 CFR Part 1303 \(Head Start privacy\)](#) · [§1302.33 screenings](#) · [§1302.17 suspension/expulsion](#) · [COPPA Final Rule 2025](#) · [IDEA confidentiality FAQ](#) · [45 CFR Part 98 \(CCDF\)](#)

STATE LAW

[CA AB 2799 / ELPIPA](#) · [NY Ed Law 2-d](#) · [Illinois BIPA](#) · [FPF state tracker](#) · [State AI legislation, mid-2026](#) · [EU AI Act Art. 5](#)

ECE & AI

[Brookings, *Generation AI Starts Early*](#) · [NCYL child-centered AI](#) · [Zero to Three, Screen Sense AI](#) · [Institute for Child Success](#) · [NAEYC position statements](#)

MODELS TO BORROW

[Council on Criminal Justice](#) · [TeachAI toolkit](#) · [Project Evident equitable AI framework](#) · [McGovern Learning Hub](#) · [Fast Forward policy builder](#) · [NTEN AI hub](#) · [CHAI applied model card](#) · [CDT K-12 policy priorities](#)

EVIDENCE BASE

[RAND pre-K teachers and AI](#) · [CSCCE Workforce Index 2024](#) · [Preschool expulsion data](#) · [MPI dual language learners](#) · [PIRG, *Trouble in Toyland 2025*](#) · [Canvas Dx real-world performance](#) · [Halberstadt et al., anger perception bias](#) · [Correction on the 2016 Yale study](#)

Working draft v2, 10 September 2026. A build spec, not a public document — nothing here has been reviewed by counsel, and the items flagged in Section 17 should be re-verified against primary sources before any of this appears on a live site.